

# Query Processing & Optimization

---

CS 377: Database Systems

# Recap: File Organization & Indexing

---

- Physical level support for data retrieval
  - File organization: ordered or sequential file to find items using binary search
  - Index: data structures to help with some query evaluation (selection & range queries)
- Indexes may not always be useful even for selection queries
- What about join queries and other queries not supported by indices?

# Query Processing Introduction

---

- Some database operations are expensive
- Performance can be improved by being “smart”
  - Clever implementation techniques for operators
  - Exploiting “equivalences” of relational operators
- Using statistics and cost models to choose better plans

# Basic Steps in Query Processing

- Parse and translate:  
convert to RA query
- Optimize RA query  
based on the different  
possible plans
- Evaluate the execution  
plan to obtain the  
query results

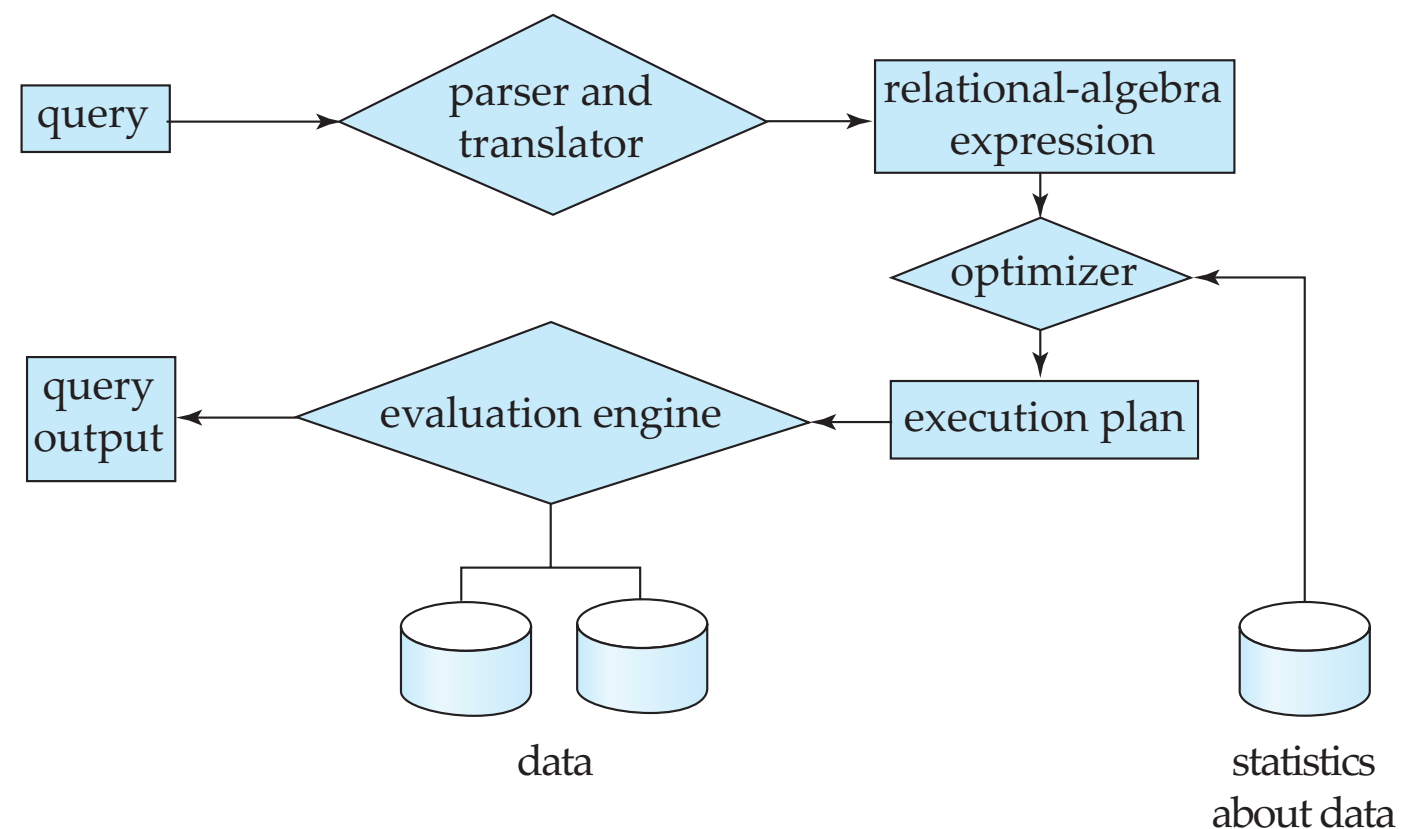


Figure 12.1 from Database System Concepts book

# Example: SQL Query

---

Find movies with stars born in 1960

```
SELECT movieTitle  
FROM StarsIn, MovieStar  
WHERE starName = name  
AND birthdate LIKE '%1960';
```

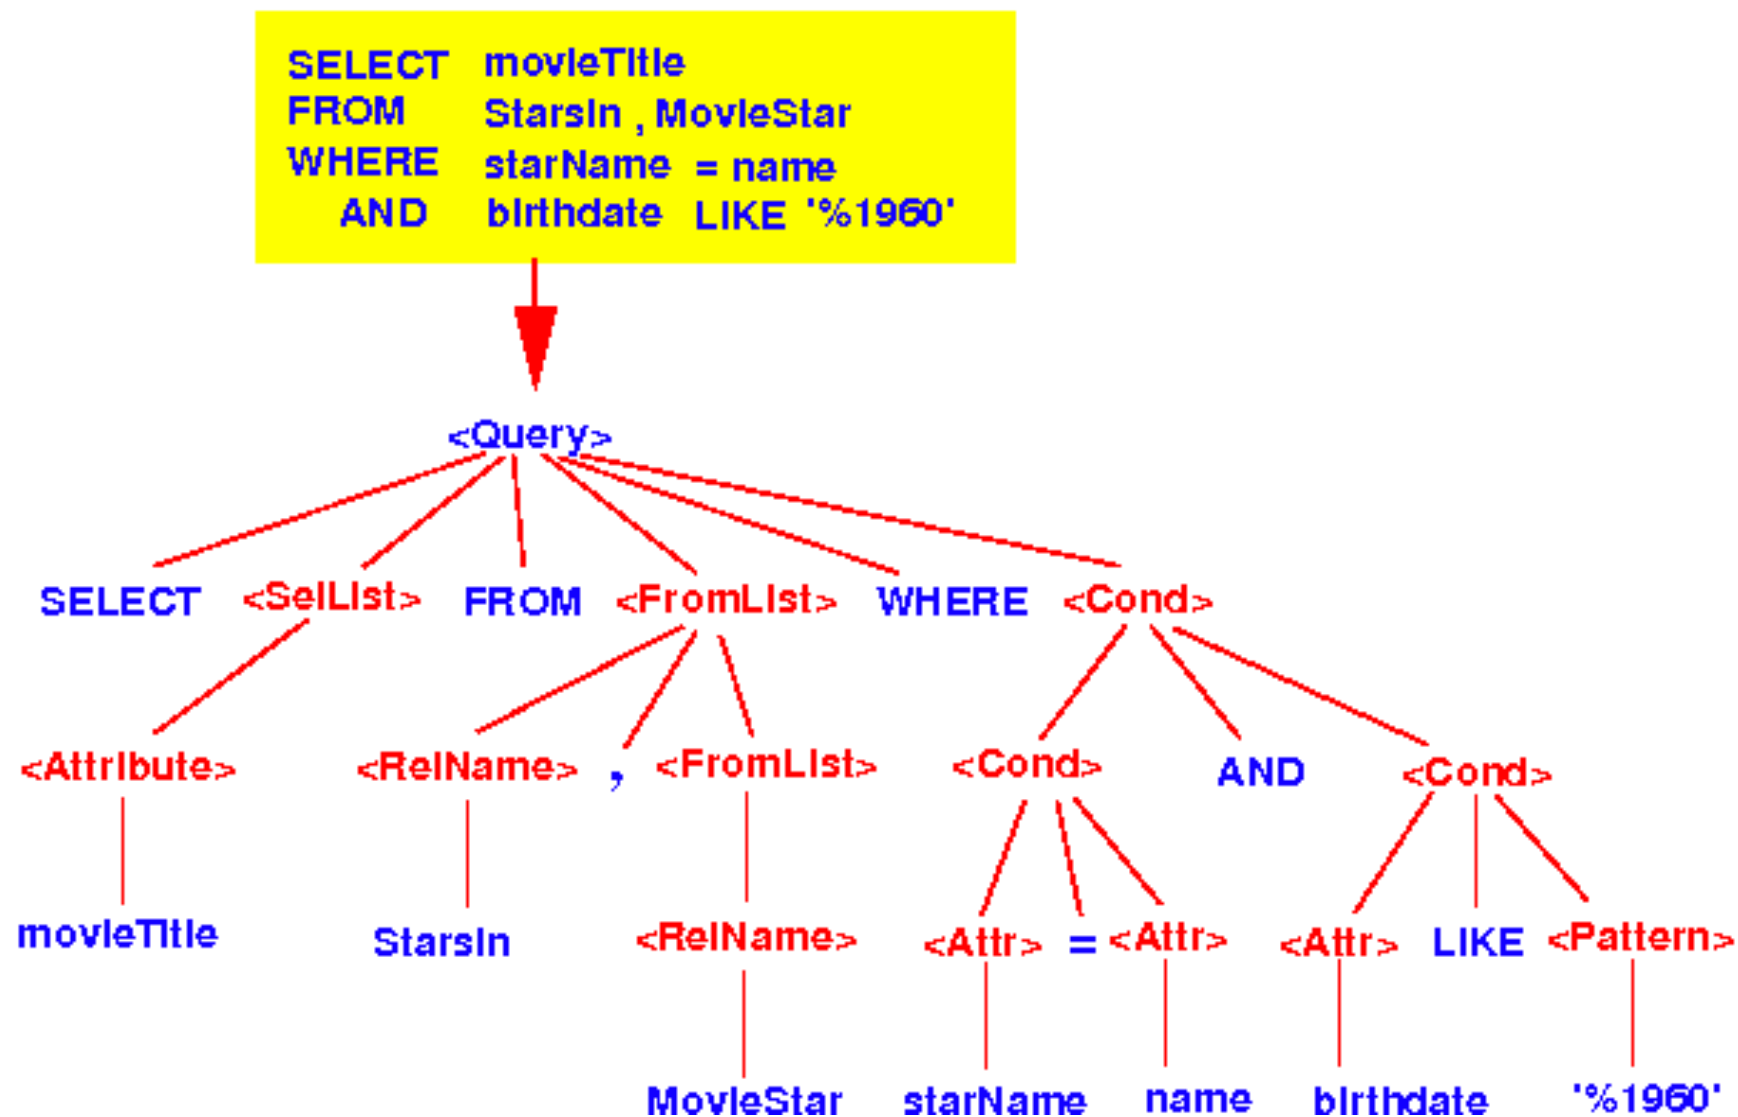
# Example: Bad Query Optimization

---

- Cartesian product first:  
StarsIn x MovieStar
  - Selection criteria next:  
starname = name AND birthdate LIKE '%1960'
  - GROUP BY; HAVING (if available)
  - Projections  
SELECT movietitle
  - ORDER BY last
- Incredibly inefficient with huge intermediate results!

# Example: SQL Query Step 1

Step 1: Convert SQL query into a parse tree

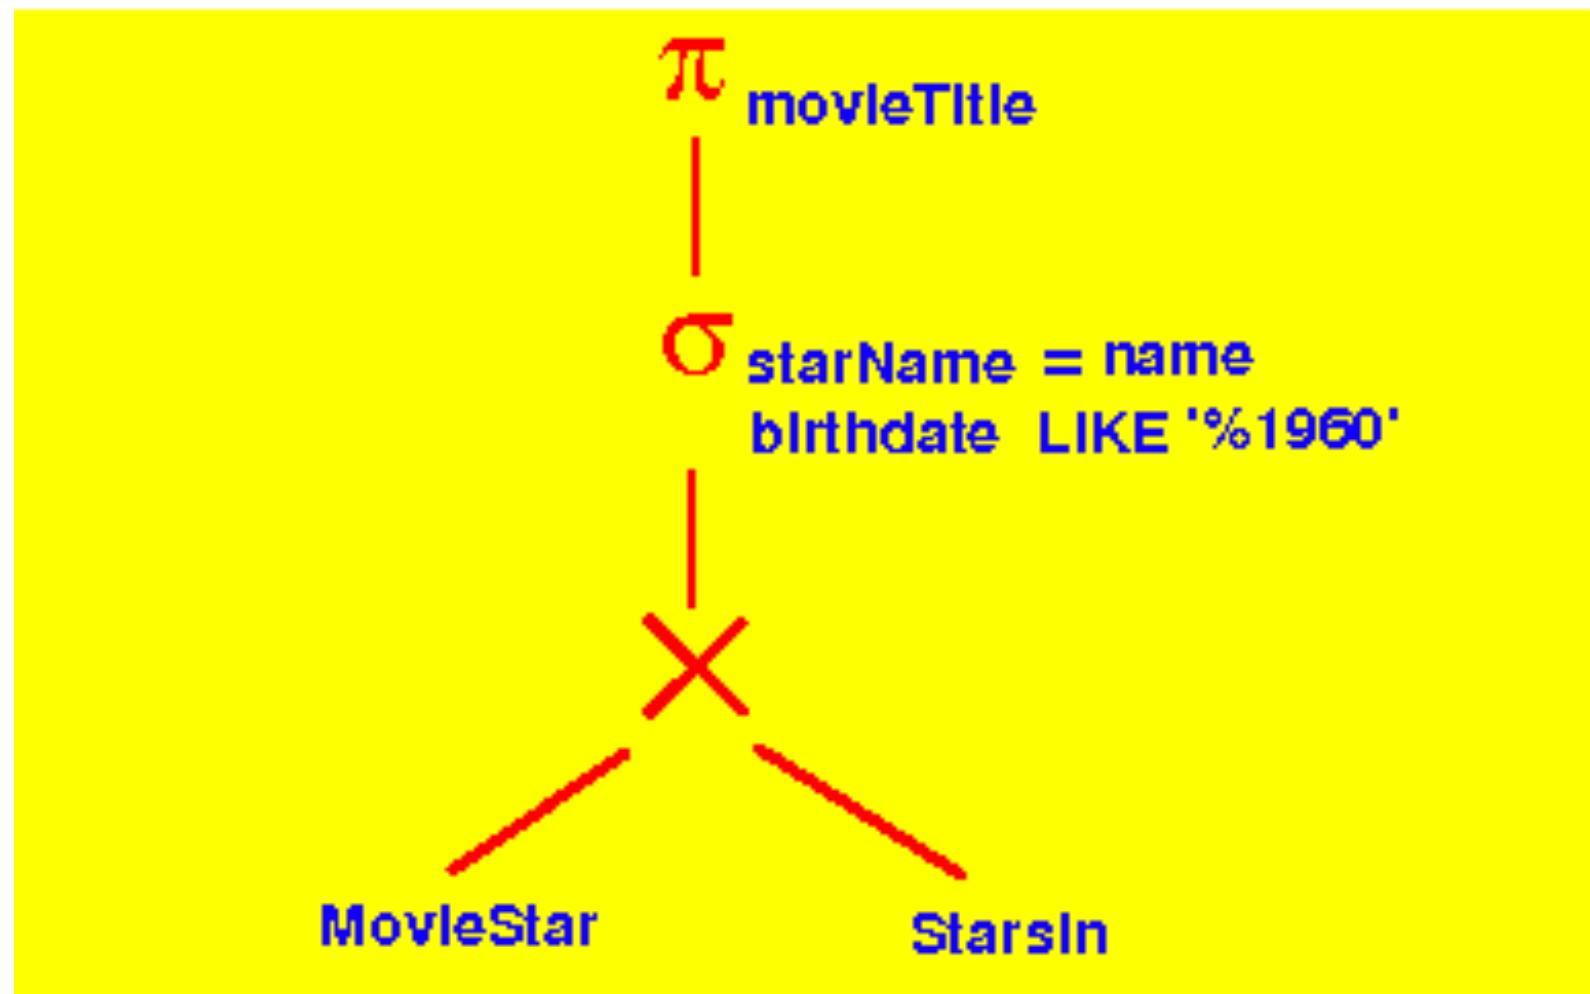


<http://www.mathcs.emory.edu/~cheung/Courses/554/Syllabus/5-query-opt/intro.html>

# Example: SQL Query Step 2

---

Step 2: Convert parse tree into initial logical query plan using RA expression

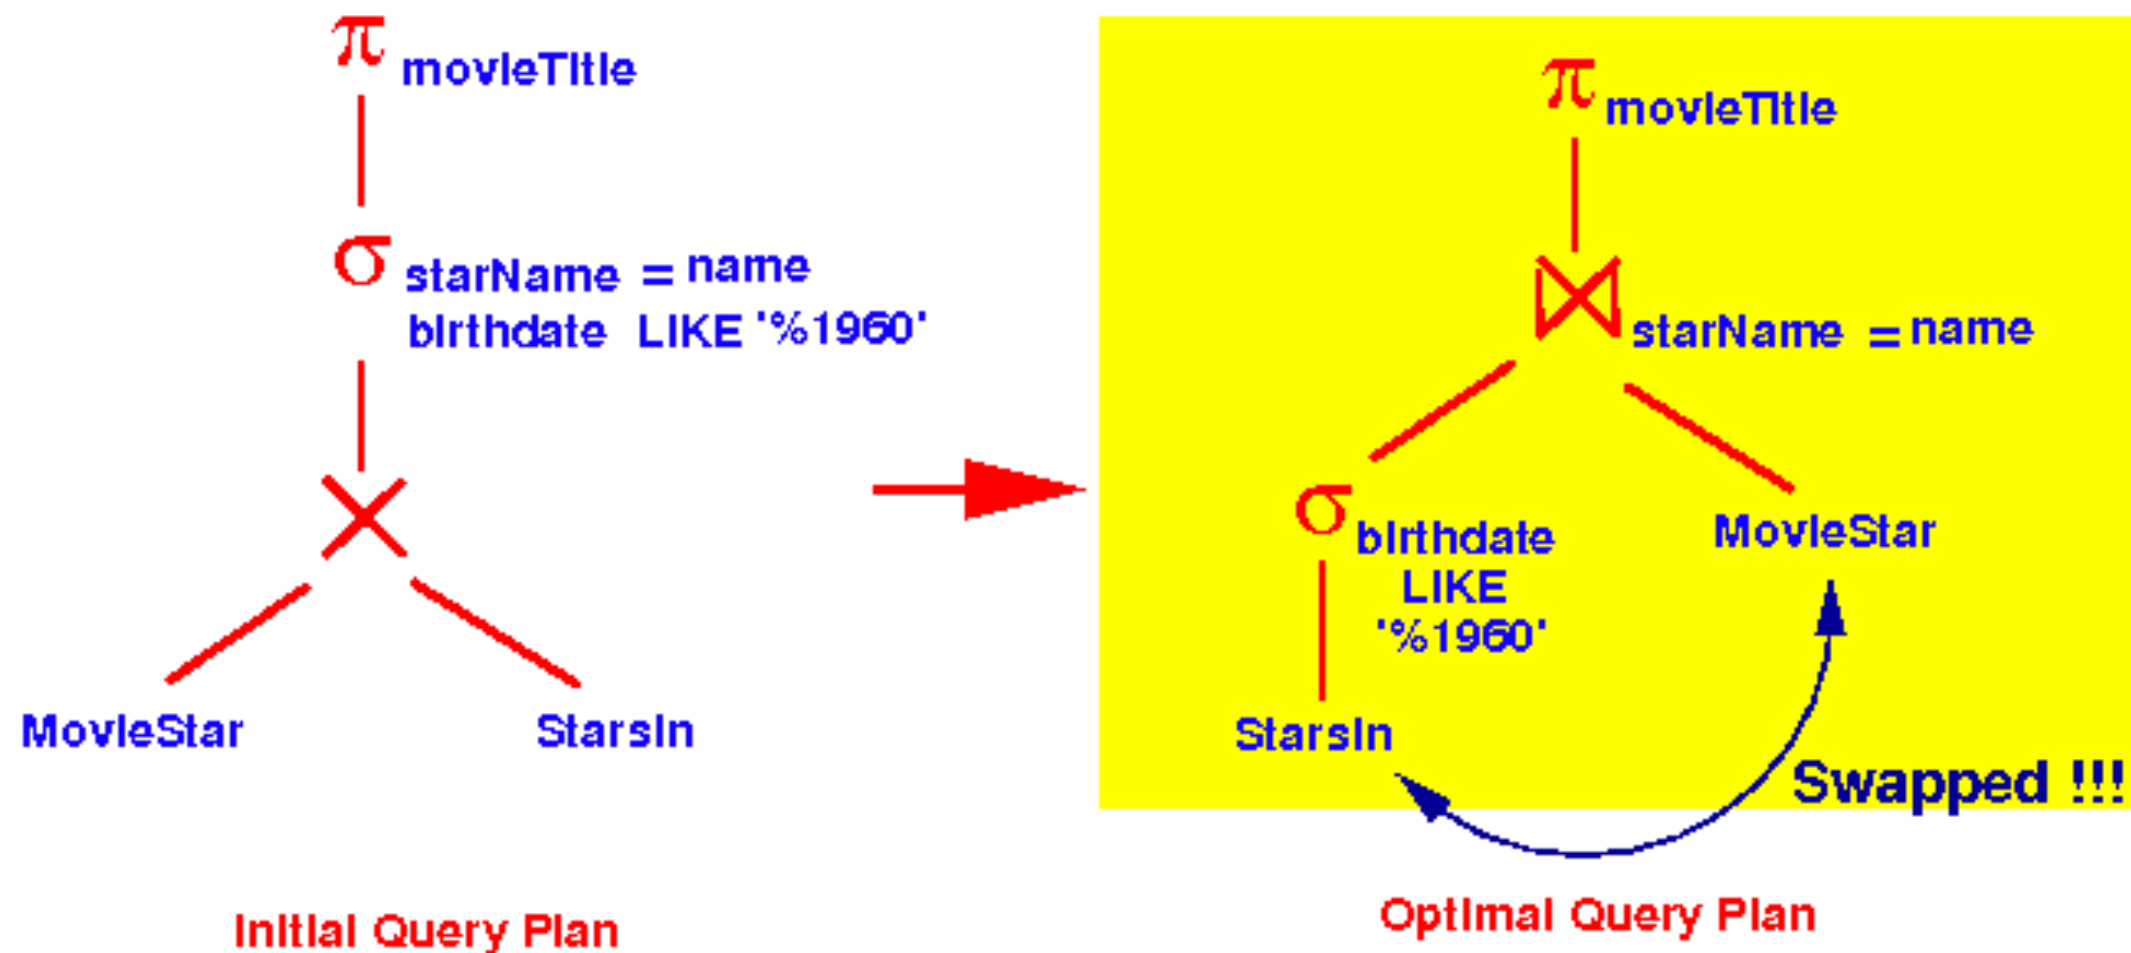


<http://www.mathcs.emory.edu/~cheung/Courses/554/Syllabus/5-query-opt/intro.html>



# Example: SQL Query Step 3

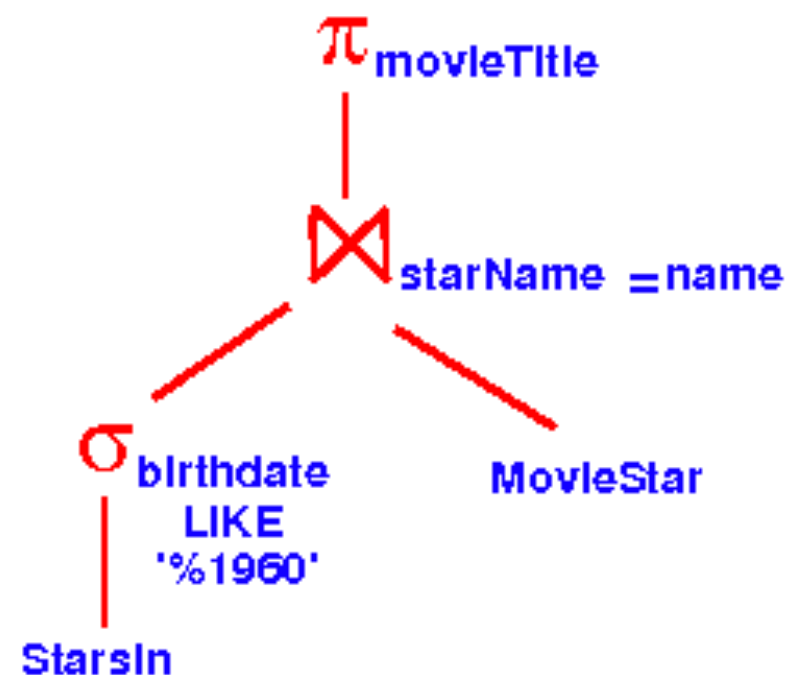
Step 3: Transform initial plan into optimal query plan using some measure of cost to determine which plan is better



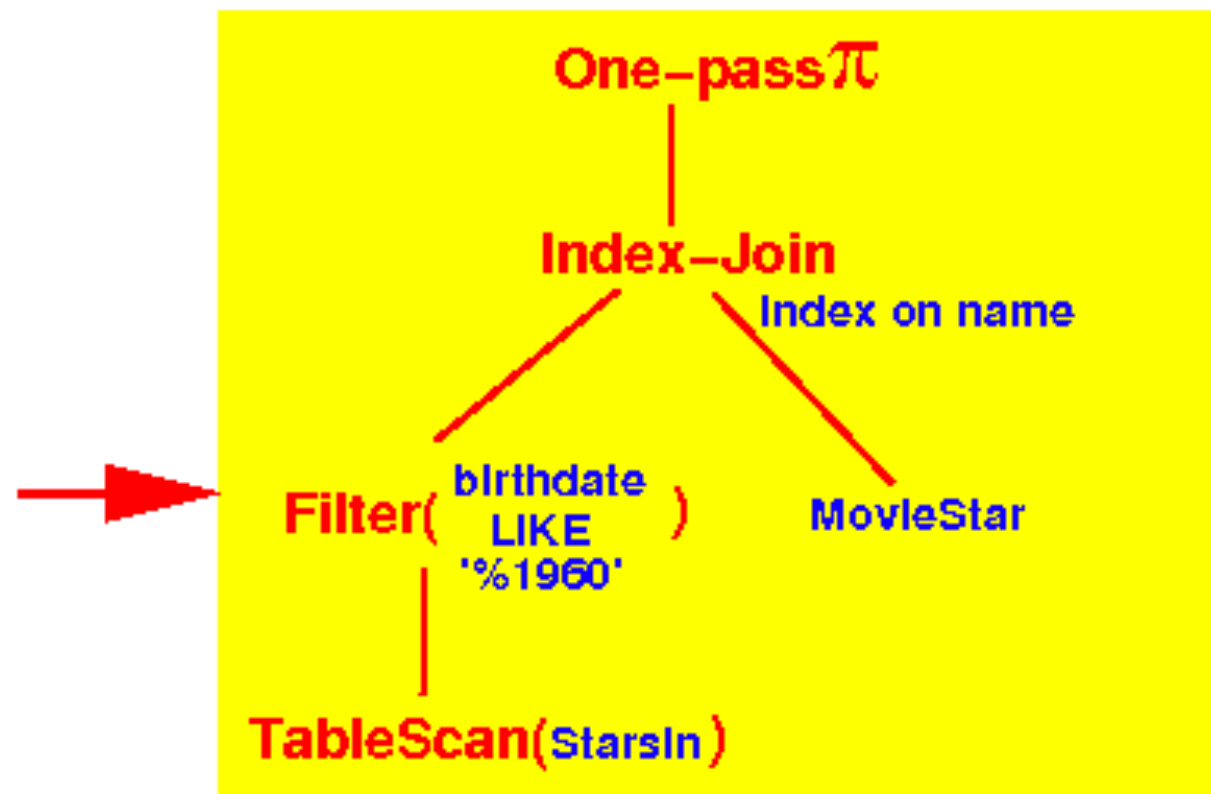
<http://www.mathcs.emory.edu/~cheung/Courses/554/Syllabus/5-query-opt/intro.html>

# Example: SQL Query Step 4

Step 4: Select physical query operator for each relational algebra operator in the optimal query plan



Optimal Logical Query Plan

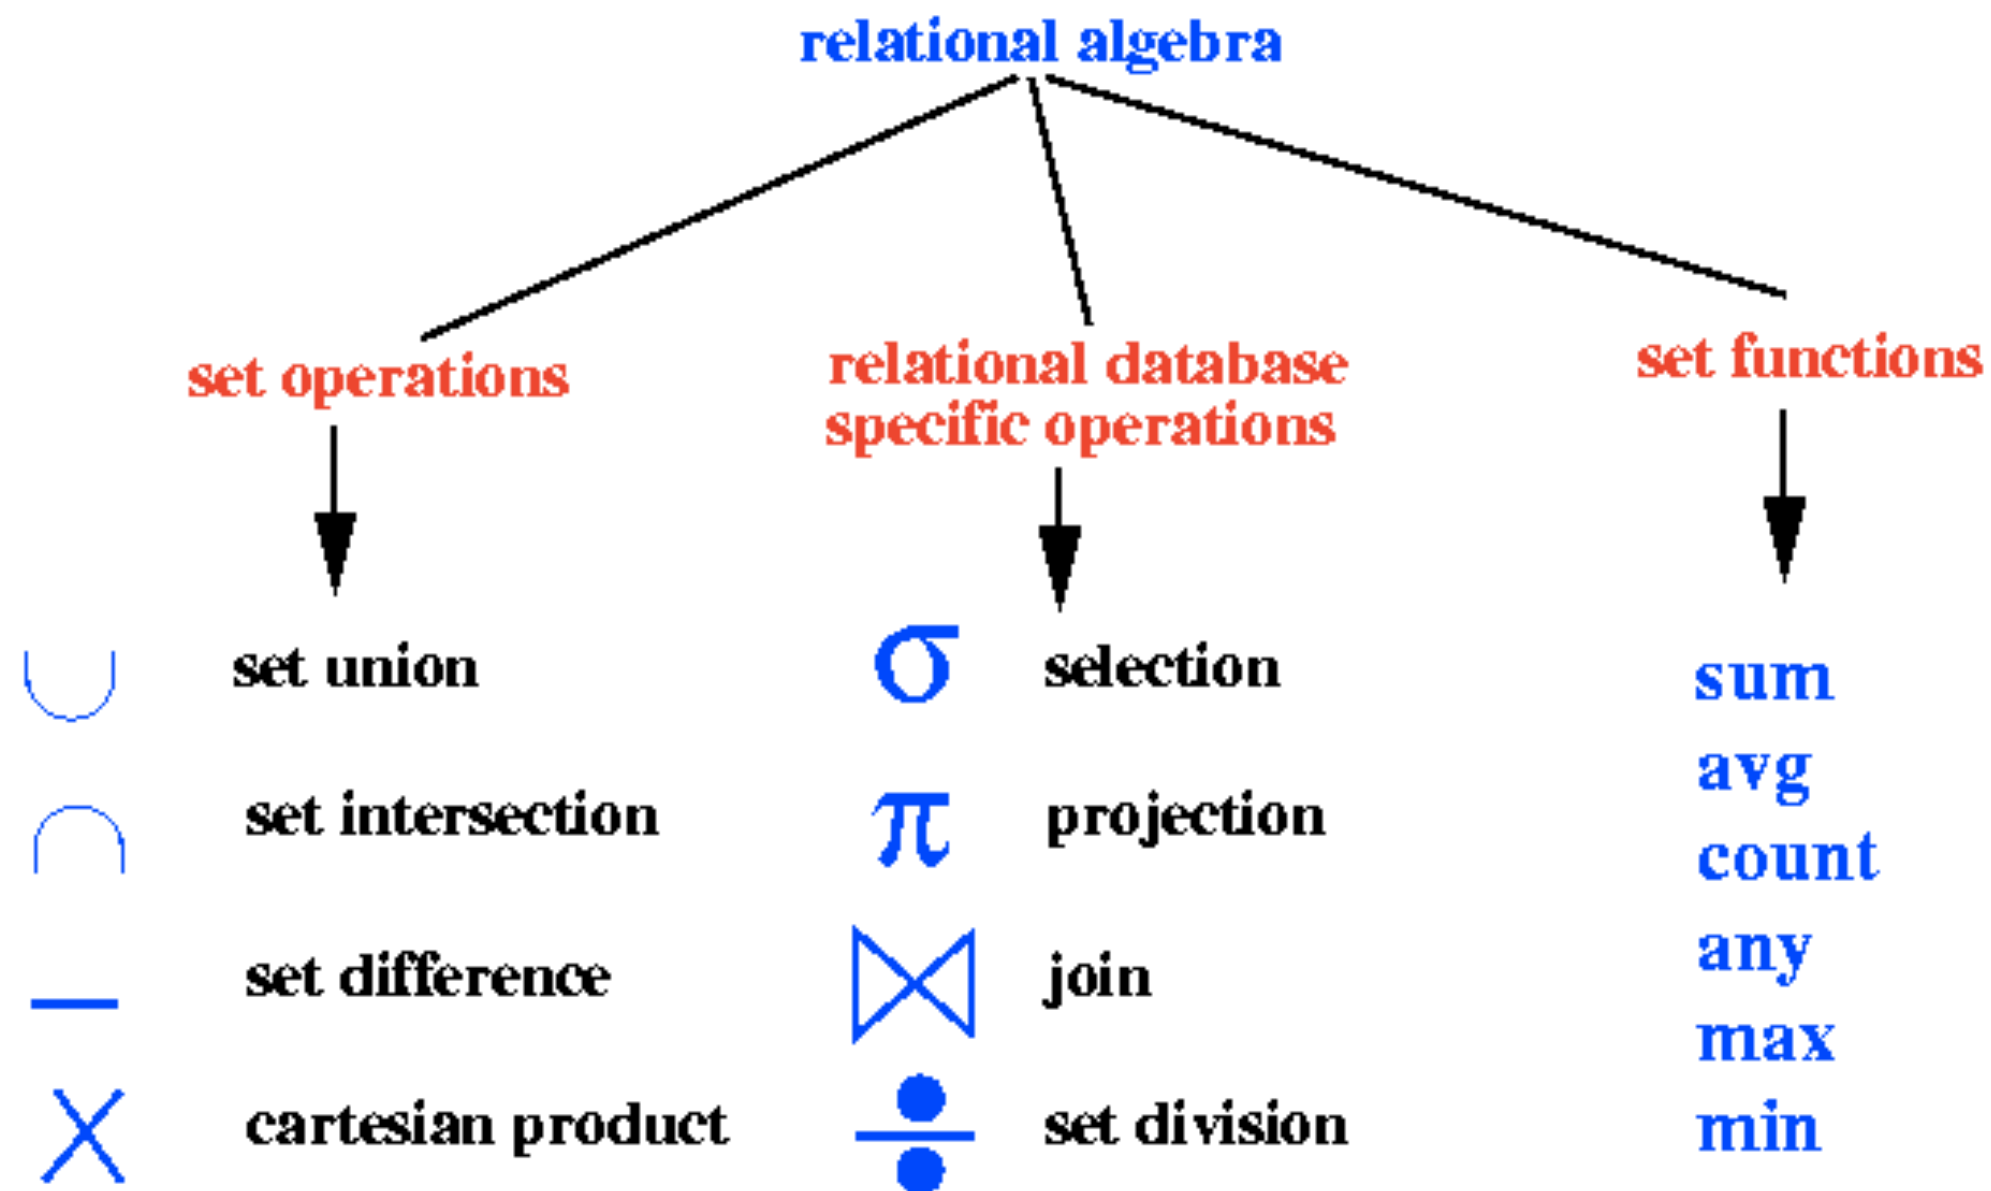


Physical Query Plan

<http://www.mathcs.emory.edu/~cheung/Courses/554/Syllabus/5-query-opt/intro.html>

# Recap: Relational Algebra

---



# Recap: SQL Query to RA


---

- How do you represent queries in RA?
- Database: Students(sid, sname, gpa)  
People(ssn, pname, address)
- SQL query:  
SELECT DISTINCT gpa, address  
FROM Students, People  
WHERE gpa > 3.5 AND sname = pname;
- RA query:  
$$\pi_{\text{gpa,address}}(\sigma_{\text{gpa} > 3.5}(\text{Students} \bowtie_{\text{sname=name}} \text{People}))$$

# Query Tree (Plan)

---

- A tree data structure that corresponds to a relational algebra expression
  - Leaf nodes = input relations
  - Internal nodes = RA operations
- Execution of query tree
  - Start at the leaf nodes
  - Execute internal node whenever its operands are available and replace node by result



# Query Optimization Heuristics

---

- Apply heuristic rules on standard initial query tree to find optimized equivalent query tree
- Main heuristic: Favor operations that reduce the size of intermediate results first
  - Apply SELECT and PROJECT operations before join or other set operations
  - Apply more selective SELECT and join first
- General transformation rules for relational algebra operators

# RA Transformation Rules

---

- Selection cascade: conjunctive selection condition can be broken into sequence of individual operations

$$\sigma_{c1 \text{ AND } c2 \text{ AND } \dots \text{ AND } c_n}(R) = \sigma_{c1}(\sigma_{c2}(\dots(\sigma_{c_n}(R))\dots))$$

- Commutativity of selection

$$\sigma_{c1}(\sigma_{c2}(R)) = \sigma_{c2}(\sigma_{c1}(R))$$

- Cascade of projection: ignore all but the last one

$$\pi_A(\pi_{A,B}(R)) = \pi_A(R)$$

- Commuting selection and projection: if the selection condition  $c$  involves only attributes in the projection list commute the two

$$\pi_{A, B}(\sigma_c(R)) = \sigma_c(\pi_{A, B}(R))$$

# RA Transformation Rules (2)

---

- Commutativity of joins, cartesian product, union, intersection

$$R \theta S = S \theta R$$

- Associativity of join, cartesian product, union, intersection

$$(R \theta S) \theta T = R \theta (S \theta T)$$

- Selection and join: if attributes in the selection condition involves only attributes of one of the relations being joined

$$\sigma_c(R \bowtie S) = \sigma_c(R) \bowtie S$$

$$\sigma_c(R \bowtie S) = \sigma_{c_1}(R) \bowtie \sigma_{c_2}(S)$$



# RA Transformation Rules (3)

---

- Commuting projection with join: if join condition involves only attributes in the projection list, commute the operations

$$\pi_L(R \bowtie_c S) = (\pi_{L_1}(R)) \bowtie_c (\pi_{L_2}(S))$$

- Commuting selection with intersection, union, or difference

$$\sigma_c(R \theta S) = (\sigma_c(R)) \theta (\sigma_c(S))$$

- Several others in the book...

# Query Optimization Heuristic Algorithm


---

- Break up any select operations with conjunctive conditions into cascade of select operations and move select operations as far down query tree as permitted
- Rearrange leaf nodes so leaf nodes with most restrictive select operations are executed first
- Combine cartesian product operation with a subsequent selection operation into join operation
- Break down and move lists of projection attributes down the tree as far as possible
- Identify subtrees that represent group of operations that can be executed as a single algorithm

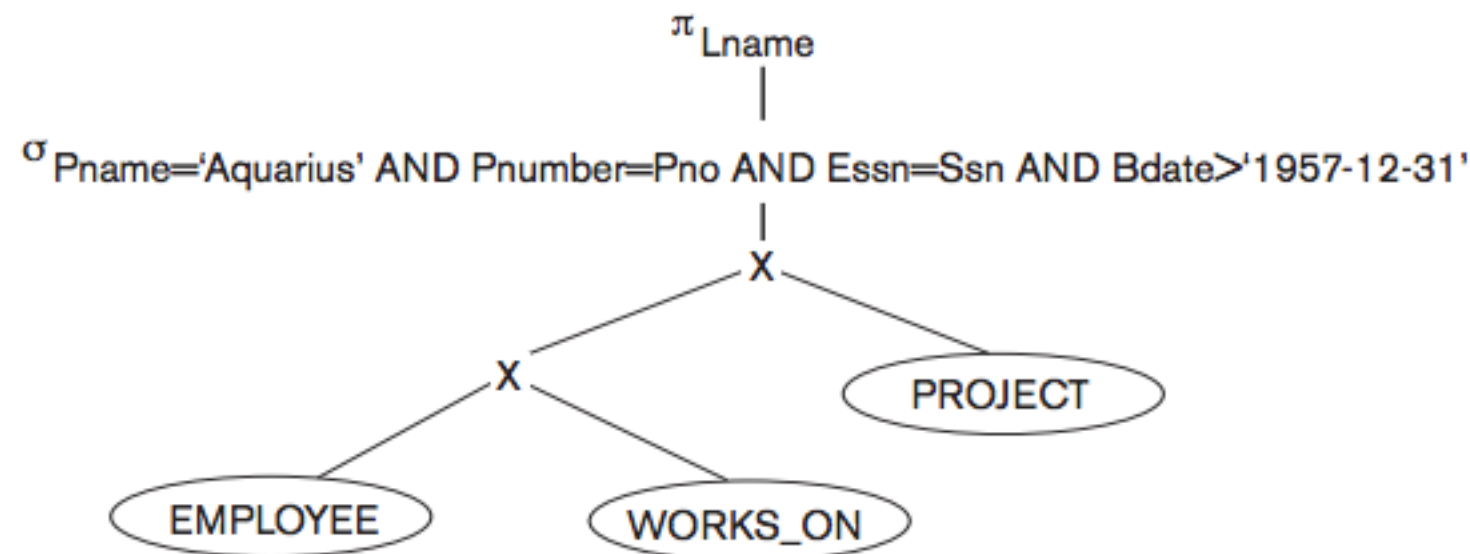
# Example: SQL Query Optimization

---

```
SELECT lname  
FROM   employee, works_on, project  
WHERE  pname = 'Aquarius' and pnumber = pno  
AND    bdate > '1957-12-31';
```

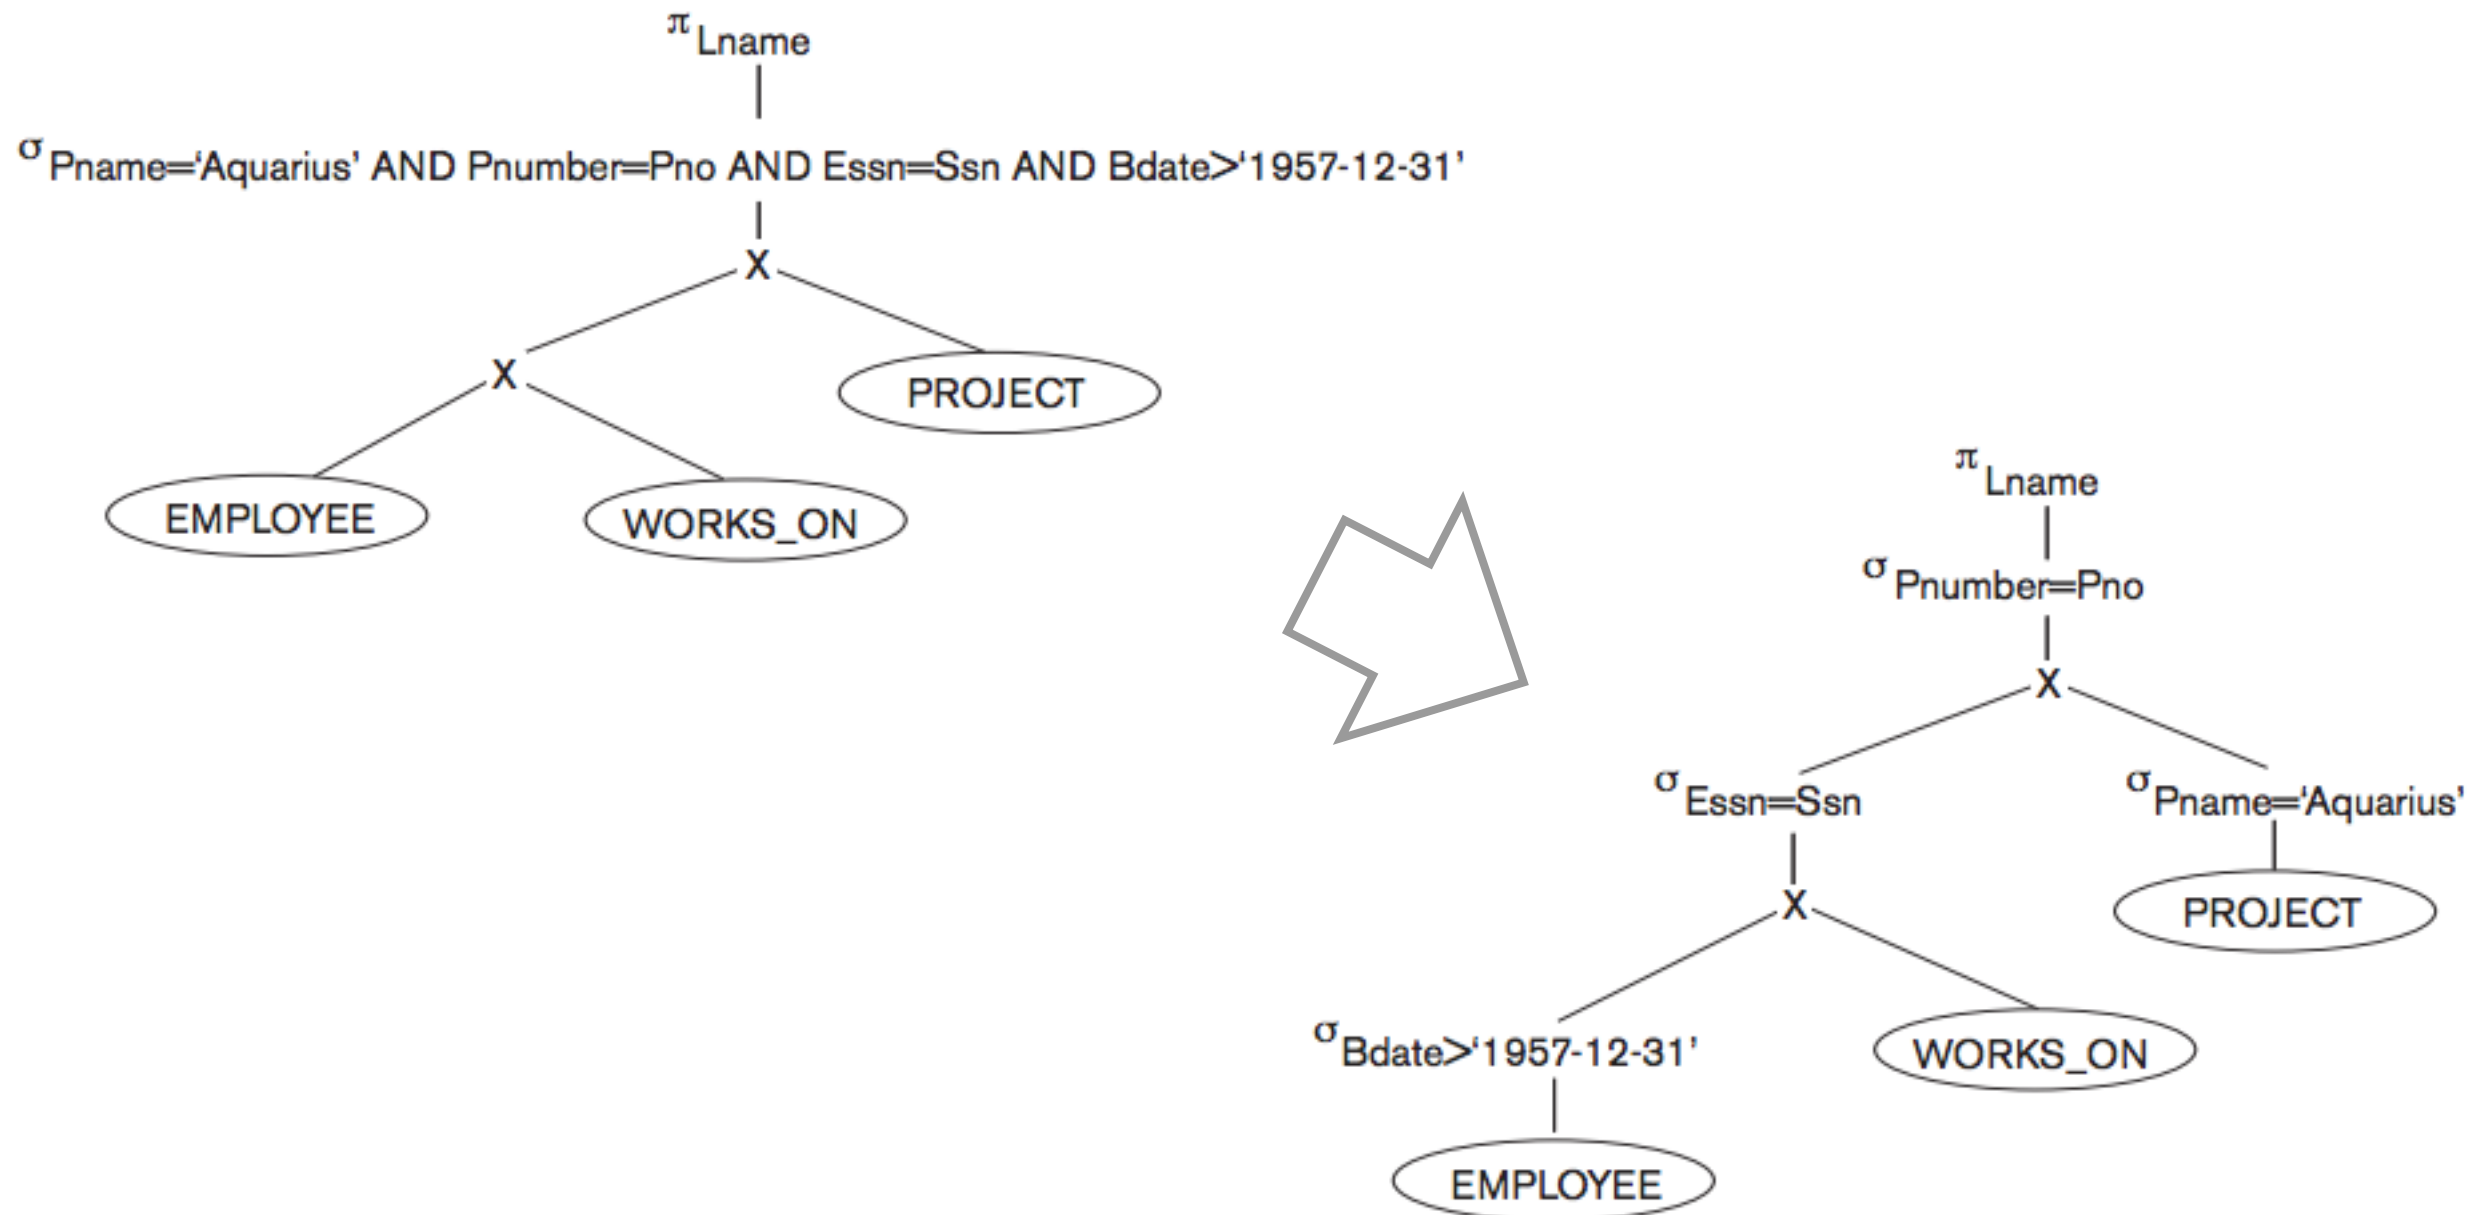


Initial query tree



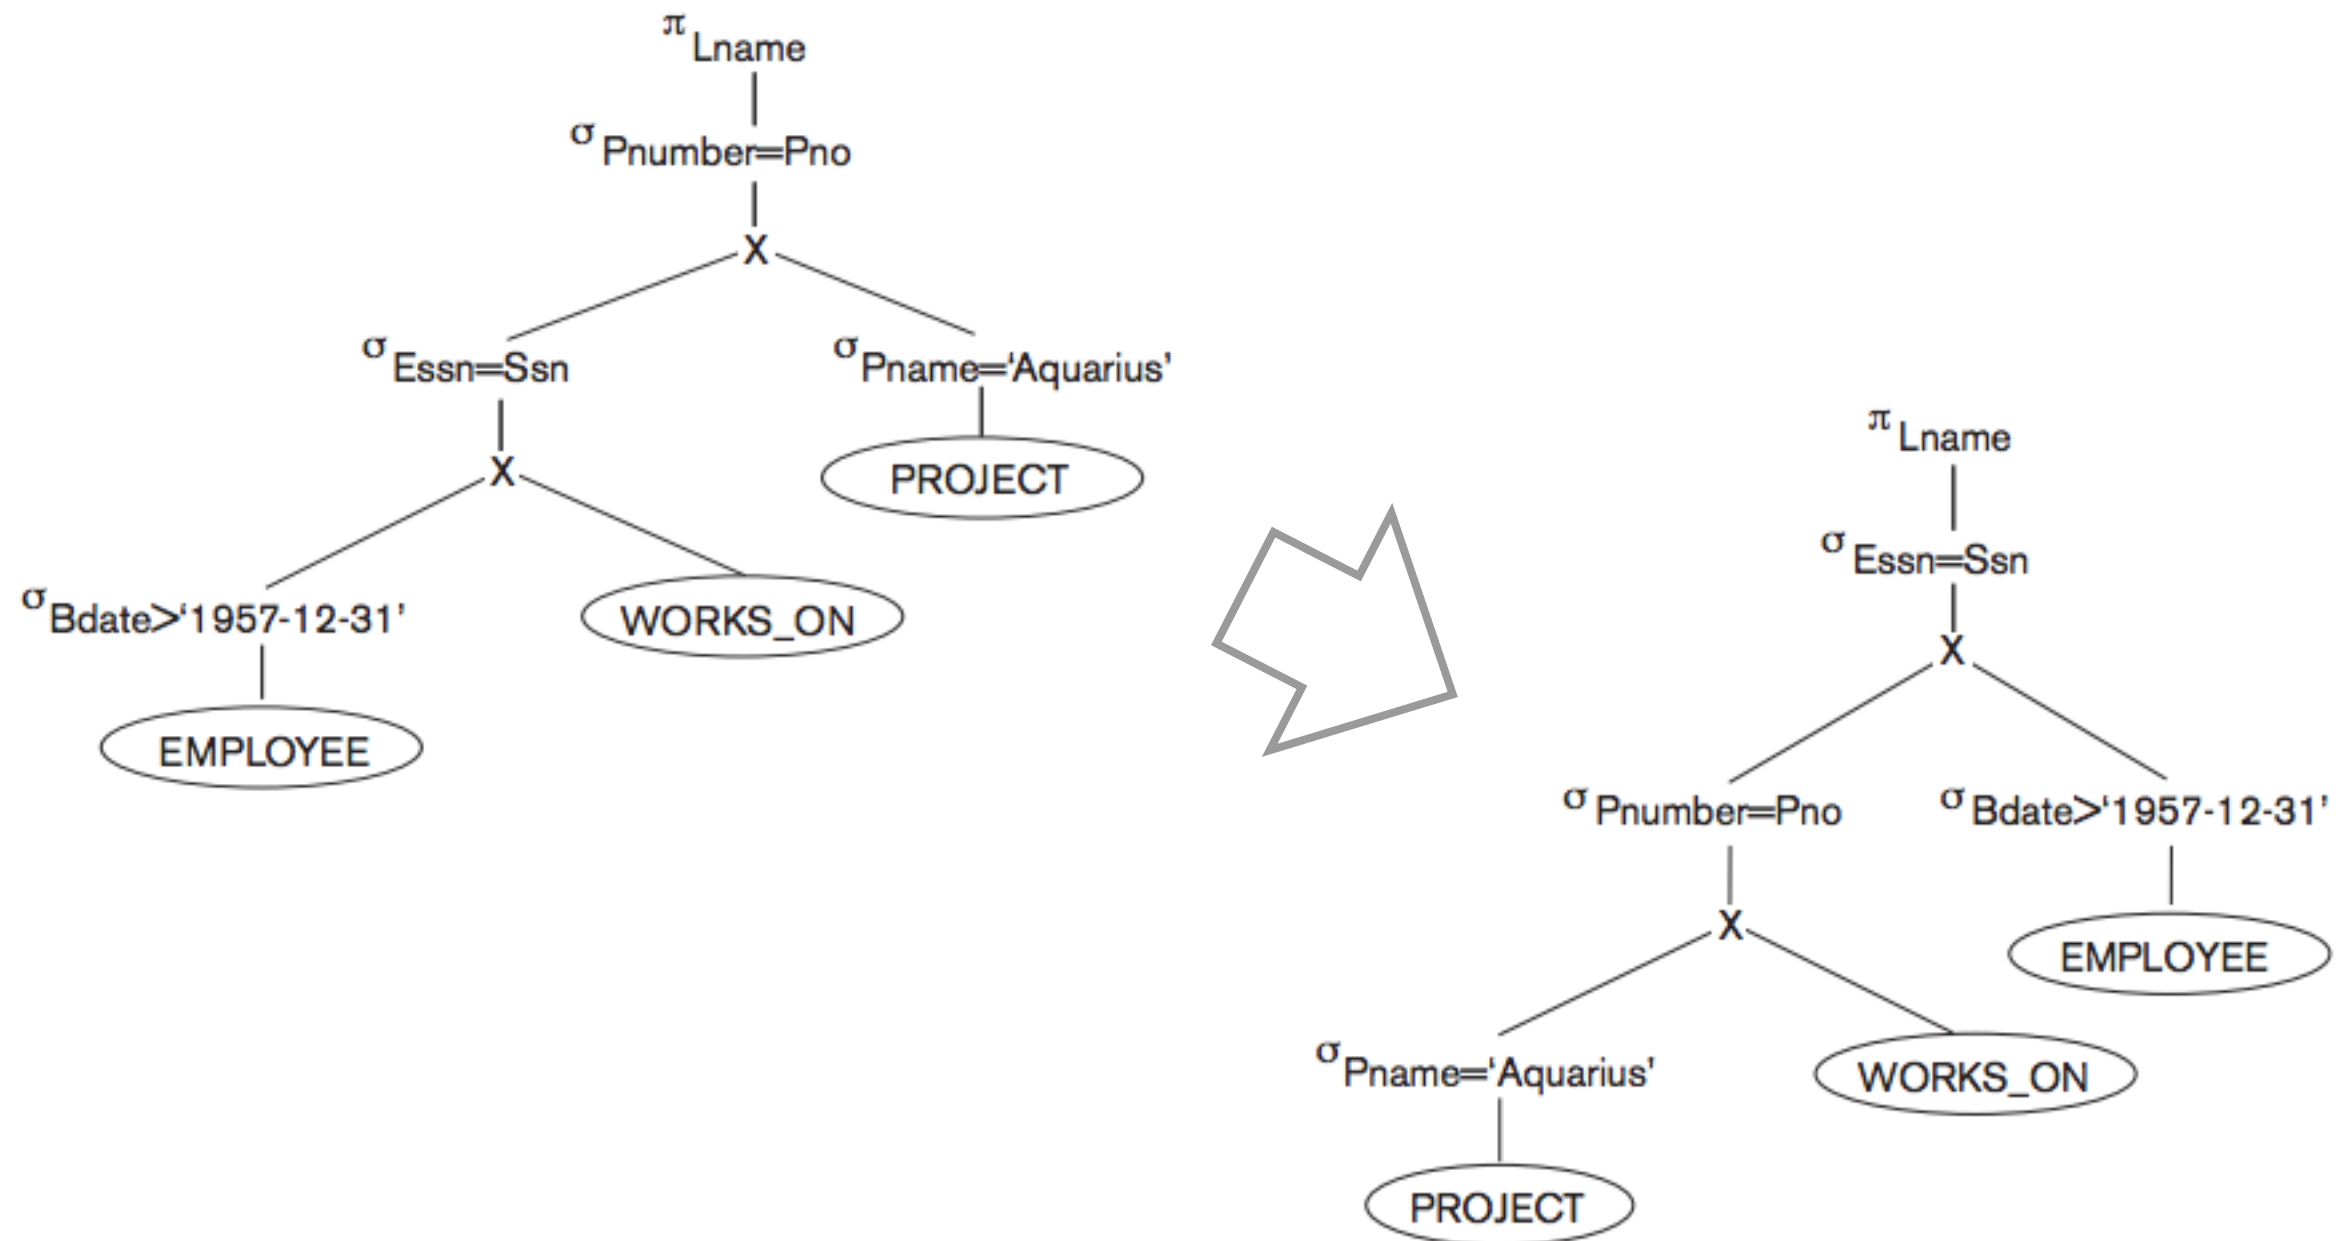
# Example: SQL Query Optimization (2)

Move SELECT operations down the query tree





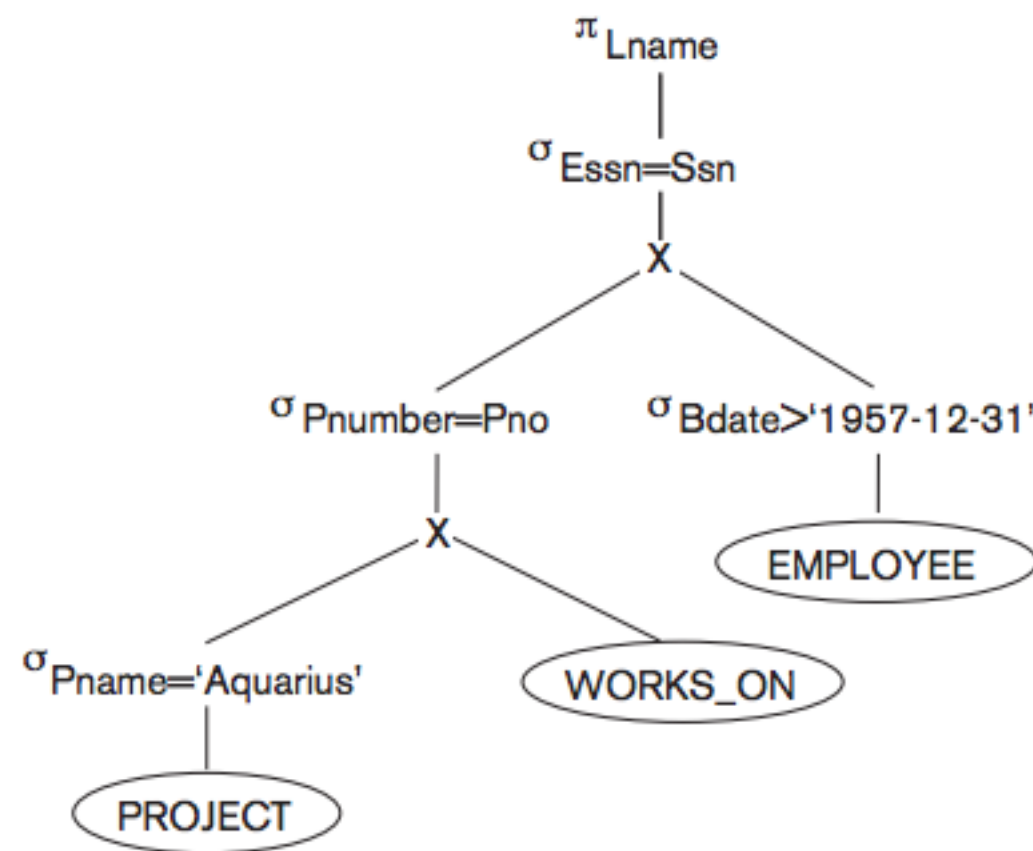
# Example: SQL Query Optimization (3)

Apply more restrictive SELECT first (left most side of tree)



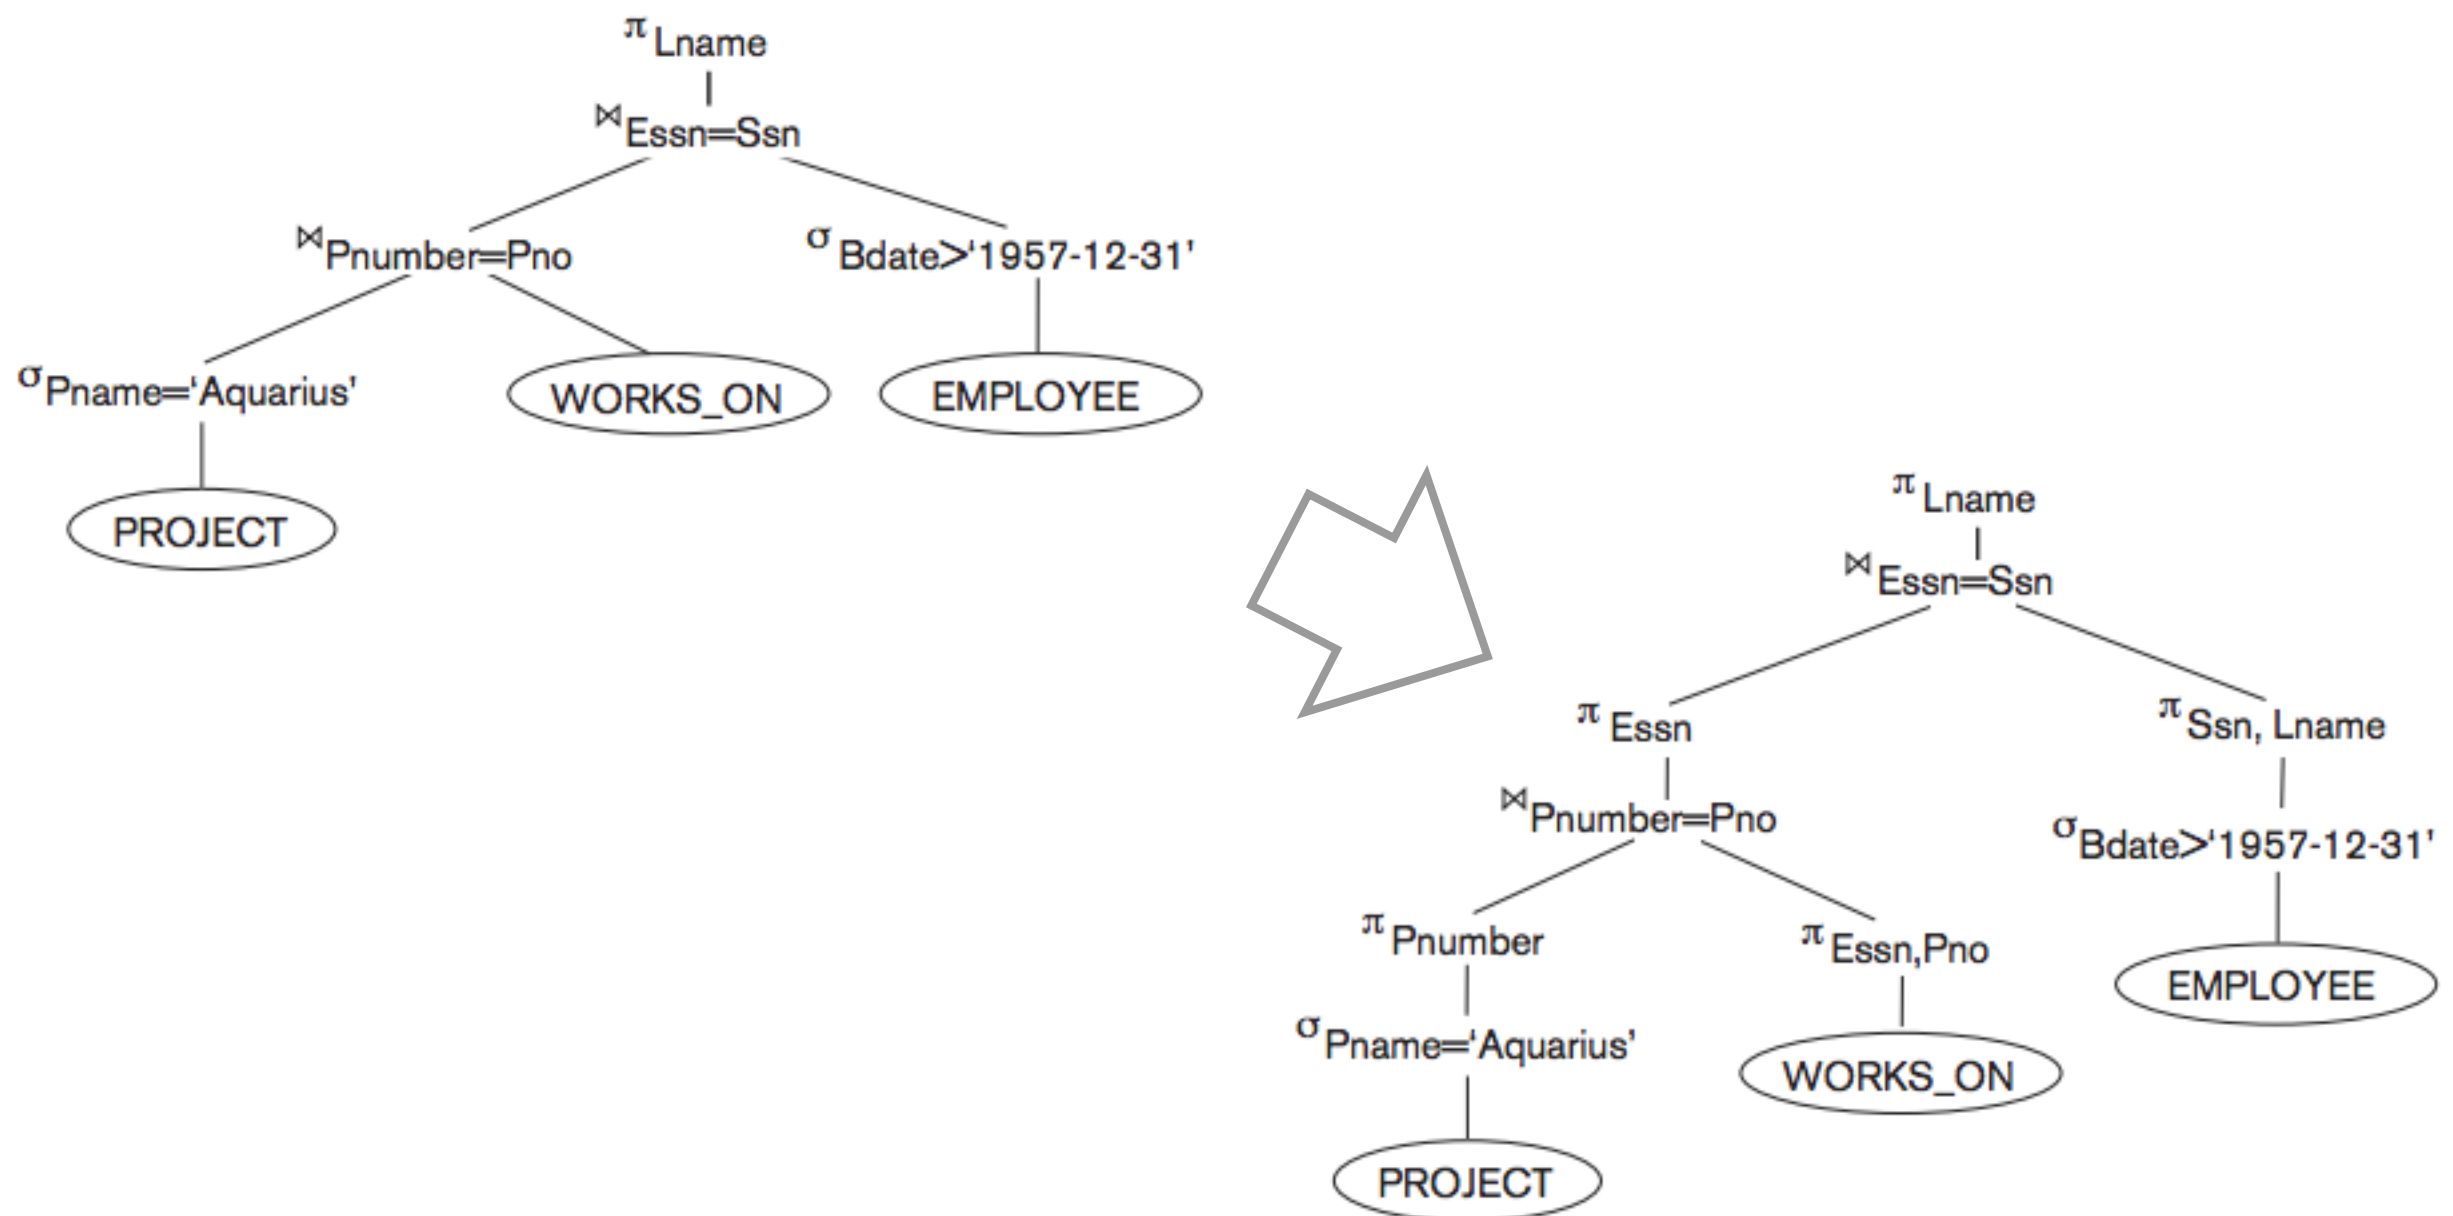
# Example: SQL Query Optimization (4)

Replace cartesian product and select with join



# Example: SQL Query Optimization (5)

Move projections down the tree



# Query Optimization

---

- Logical level: heuristics based optimization to find a better RA query tree
- SQL query  $\rightarrow$  initial logical query tree  $\rightarrow$  optimized query tree
- Physical level: cost-based optimization to determine “best” query plan
- Optimized query tree  $\rightarrow$  query execution plans  $\rightarrow$  cost estimation  $\rightarrow$  “best” query plan



# Cost-based Query Optimization

---

Estimate and compare the costs of executing a query using different execution strategies and choose the strategy with the lowest cost estimate

- Disk I/O cost
- Storage cost
- Computation cost
- Memory usage cost
- Communication cost (distributed databases)

# Catalog Information

---

Database maintains statistics about each relation

- Size of file: number of tuples  $[n_r]$ , number of blocks  $[b_r]$ , tuple size  $[s_r]$ , number of tuples or records per block  $[f_r]$ , etc.
- Information about indexes and indexing attributes
  - Attribute values - number of distinct values  $[V(\text{att}, r)]$
  - Selection cardinality - expected size of selection given value  $[SC(\text{att}, r)]$
  - ...

# Catalog Information for Index

---

- Average fan-out of internal nodes of index  $i$  for tree-structured indices  $[f_i]$
- Number of levels in index  $i$  (i.e., height of index  $i$ )  $[HT_i]$ 
  - Balanced tree on attribute  $A$  of relation  $r$ :  $\lceil \log_{f_i} V(A, r) \rceil$
  - Hash index: 1
- Number of lowest-level index blocks in  $i$  (i.e., number of blocks at the leaf level of the index)  $[LB_i]$

# Example: Bank Schema

---

## Account relation

- $f_{\text{account}} = 20$  (20 tuples per block)
- $V(\text{bname}, \text{account}) = 50$  (50 branches)
- $V(\text{balance}, \text{account}) = 500$  (500 different balance values)
- $n_{\text{account}} = 10000$  (10,000 tuples in account)
- $b_{\text{account}} = 10000 / 20 = 500$

# SELECT Algorithms (Simple)

---

- Linear search (brute force): selection attribute is not ordered and no index on attribute
  - Cost: # blocks in relation =  $b_r$
  - Reserves example: 500 I/Os
- Binary search: selection attribute is ordered and no index
  - Cost:  $\underbrace{\lceil \log_2(b_r) \rceil}_{\text{locating first tuple}} + \underbrace{\lceil \text{SC}(\text{att}, r) / f_r \rceil}_{\text{\# blocks with selection}} - 1$

# Example: Binary search

---


- How expensive is the following query if we assume Account is sorted by branch name?

$$\sigma_{\text{bname}=\text{'Perryridge'}}(\text{Account})$$

- Ans:
  - # of tuples in the relation pertaining to Perryridge is total number of tuples divided by distinct values:  $10000/50$
  - Cost:  $\lceil \log_2(500) \rceil + \lceil 200/20 \rceil - 1 = 18$

# SELECT Algorithms (Simple w/ Index)

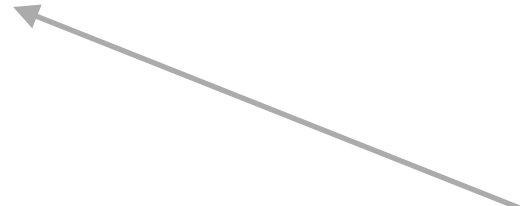
---

- Index search: cost depends on the number of qualifying tuples, cost of retrieving the tuples and the type of query
  - Primary index
    - Equality search on candidate key:  $HT_i + 1$
    - Equality search on nonkey:  $HT_i + \lceil SC(\text{att}, r) / f_r \rceil$
    - Comparison search:  $HT_i + \lceil c / f_r \rceil$   
  
estimated number of tuples that satisfy condition

# SELECT Algorithms (Simple w/ Index)

---

- Secondary index
  - Equality search on candidate key:  $HT_i + 1$
  - Equality search on nonkey:  $HT_i + SC(att, r)$
  - Comparison search:  $HT_i + LB_i * c / n_r + c$



Note that linear file scan maybe cheaper if the number of tuples satisfying the condition is large!



# Example: Index search

---

- How expensive is the following query if we assume primary index on branch name?

$$\sigma_{\text{bname}=\text{'Perryridge'}}(\text{Account})$$

- Ans:
  - 200 tuples relating to Perryridge branch => clustered index
  - Assume B<sup>+</sup>-tree index stores 20 pointers per node, then index must have between 3 and 5 leaf nodes with a depth of 2
  - Cost:  $2 + \lceil 200/20 \rceil = 12$

# SELECT Algorithms (Complex)

---

- Conjunctive selection (several conditions with AND)
  - Single index: retrieve records satisfying some attribute condition (with index) and check remaining conditions
  - Composite index
  - Intersection of multiple indexes
- Disjunctive selection (several conditions with OR)
  - Index/binary search if all conditions have access path and take union
  - Linear search otherwise

# Example: Complex search

---

- How expensive if we want to find accounts where the branch name is Perryridge with a balance of 1200 if we assume there is a primary index on branch name and secondary on balance?
- Ans for using one index:
  - Cost for branch name: 12 block reads
  - Balance index is not clustered, so expected selection is  $10,000 / 500 = 20$  accounts
  - Cost for balance:  $2 + 20 = 22$  block reads
  - Thus use branch name index, even if it is less selective!

# Example: Complex search (2)

---

- Ans for using intersection of two indexes:
  - Use index on balance to retrieve set of S1 pointers: 2 reads
  - Use index on branch name to retrieve set of S2 pointers: 2 reads
  - Take intersection of the two
  - Estimate 1 tuple in  $50 * 500$  meets both conditions, so we estimate the intersection of two has one pointer
  - Estimated cost: 5 block reads

# Sorting

---

- One of the primary algorithms used for query processing
  - ORDER BY
  - DISTINCT
  - JOIN
- Relations that fit in memory — use techniques like quicksort, merge sort, bubble sort
- Relations that don't fit in memory — external sort-merge

# External Sort-Merge

---

- Problem: Sort  $r$  records, stored in  $b$  file blocks with a total memory space of  $M$  blocks
- Create sorted runs with  $i = 0$ 
  - Read  $M$  blocks of relation into memory
  - Sort the in-memory blocks
  - Write sorted data to run  $R_i$ , increment  $i$

# External Sort-Merge (2)

---

- Merge the sorted runs: merge subfiles until 1 remains
  - Select the first record in sort order from each of the buffers
  - Write the record to the output
  - Delete the record from the buffer page, and read the next block if empty
- Total cost:  $b_r(2\lceil \log_{M-1}(b_r/M) \rceil + 1)$

# Example: External Merge Sort

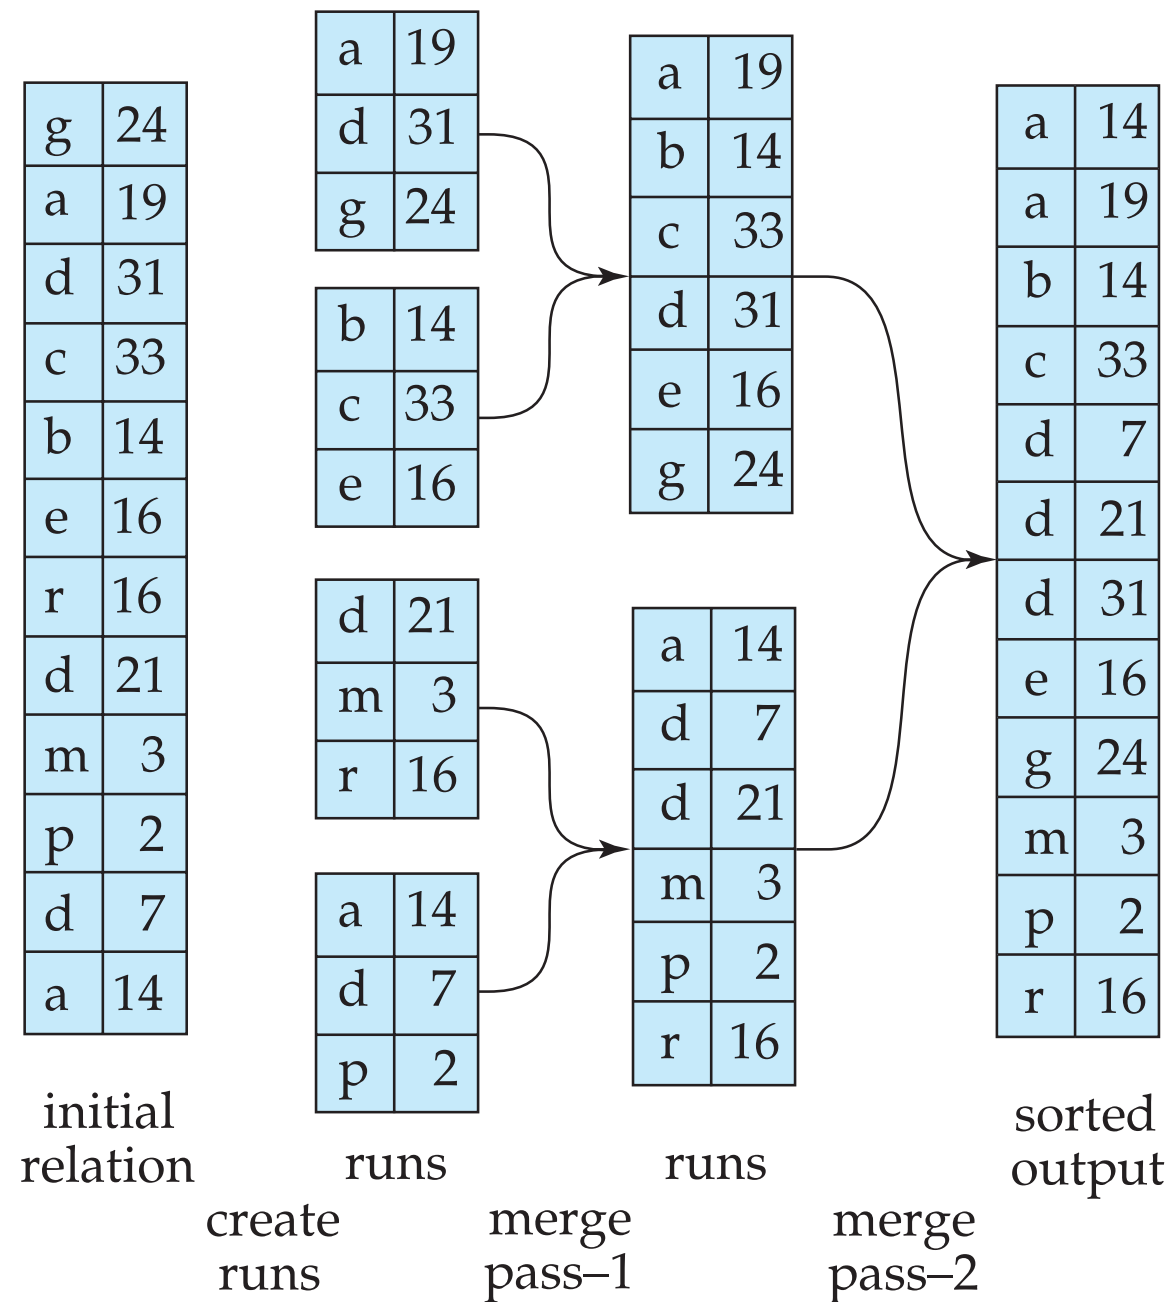


Figure 12.4 from Database System Concepts book



# Query Processing & Optimization: Recap

---

- Motivation for query optimization
- Query parse tree
- Query optimization heuristics
  - RA transformation rules
- Cost-based query optimization
  - SELECT
  - Sorting

