

SQL: Data Update & View Definition

CS 377: Database Systems

Recap: SQL Queries

SELECT [**DISTINCT**] <attribute list>
FROM <table list>
[WHERE <condition on the tables>]
[GROUP BY <grouping attributes>]
[HAVING <group condition>]
[ORDER BY <attribute list> **ASC | DESC**]
[LIMIT <number of tuples>]

SQL Query: Temporal Relation

- Result of a **SELECT** clause that exists temporally, which assists you in formulating a query
- Syntax:
SELECT <attributes>
FROM R1, R2, (SELECT ...) <alias>, ..., RN
WHERE <condition>;
- Must always use an alias to denote the result relation of the **SELECT** command

SQL Example: Temporal Relation

Find fname, lname of male employees with salary > 50K

```
SELECT *  
FROM    (SELECT fname, lname, salary  
          FROM    employee  
          WHERE sex = 'M') r1  
WHERE   r1.salary > 50000
```

SQL Query: Temporal Relation Notes

- You can use multiple temporal relations
- You cannot use a temporal relation to create another temporal relation

Example of incorrect usage:

```
SELECT ...  
FROM    ...,  
        (SELECT ...) r1,  
        (SELECT ... FROM r1 ...) r2,  
WHERE ...;
```

SQL Query: WITH

- SQL-99 standard introduced **WITH** clause to help refine the result of a query (another way to achieve temporal relation)
 - Some vendors do not support **WITH** (e.g., MySQL)
- Syntax:
WITH <alias> AS (SELECT ...)[,
 <alias2> AS (SELECT ...)]
SELECT <query>;
- Can be used to perform “refinement” on a query
 - Subsequent queries in the WITH clause can use the results of the previous query

SQL Example: WITH

Find all information on dependents of John Smith

```
WITH r1 as (SELECT *  
             FROM employee  
             WHERE fname = 'John'  
                   AND lname = 'Smith')  
  
SELECT *  
FROM   dependent  
WHERE  essn IN (SELECT ssn from r1);
```

SQL Query: JOIN Operations

- SQL-99 standard added several join operations:
 - **INNER JOIN** (normal join)
 - **LEFT JOIN** (left outer join)
 - **RIGHT JOIN** (right outer join)
 - **FULL JOIN** (outer join)
- Each operation results in a relation
- Operation can only appear in:
 - **FROM** clause of **SELECT** command
 - **WHERE** clause of **SELECT** command with an operator that uses a sub-query

SQL Query: [INNER] JOIN

- Compute the (inner) join between tables r_1 and r_2 with a given join condition
- Syntax:
 r_1 JOIN r_2 ON <join-condition>;
or
 r_1 INNER JOIN r_2 ON <join-condition>;
- JOIN operator makes the SQL query look a lot like RA query
- Can join more than 2 relations

SQL Example: INNER JOIN

Find fname, lname of employees in the 'Research' department

RA Query:

$\pi_{\text{fname}, \text{lname}}(\sigma_{\text{dname}=\text{'Research'}}(\text{EMPLOYEE} \bowtie_{\text{dno}=\text{dnumber}} \text{DEPARTMENT}))$

SQL Query:

```
SELECT fname, lname
FROM    (employee JOIN department
           ON dno = dnumber)
WHERE  dname = 'Research';
```

SQL Query: OUTER JOIN

- Compute the outer join between tables r_1 and r_2 with a given join condition - see RA slides for details on difference between left, right, and full outer joins
- Syntax:
 r_1 LEFT | RIGHT | FULL [OUTER] JOIN r_2 on <join condition>;
- Results in NULL values for the attributes where non-matching tuples occur

SQL Query: NATURAL JOIN

- Compute the natural join on attributes with the same names from two or more tables with the common attribute appearing only once in the result
- Syntax:
r1 NATURAL JOIN r2;
- Example:
SELECT *
FROM works_on NATURAL JOIN dependent;

SQL Query: CROSS JOIN

- Cross join is the same as a Cartesian Product
- Syntax:
r1 CROSS JOIN r2;
- Example:
SELECT ssn, fname, lname, dno, dnumber, dname
FROM employee CROSS JOIN dependent;

SQL Outline

- Data definition
 - Database Creation
 - Table Creation
- Query (SELECT)
- Data update (INSERT, DELETE, UPDATE)
- View definition



SQL

SQL Modifications/Updates

- A modification command does not return a result but it changes the database
- There are 3 kinds of modifications
 - **INSERT** tuple(s)
 - **DELETE** tuple(s)
 - **UPDATE** the value(s) of existing tuples

SQL Modification: INSERT

- Add one more more tuples to an existing relation
- Two forms of **INSERT**:
 - Literal values (constant or known values)
 - Result from a **SELECT** command

SQL Modification: INSERT (2)

Inserting a tuple using literal/constant values

Syntax:

INSERT INTO <table name>[(<attr names>)]
VALUES (<list of values>);

- Complete tuple: omitting [(<attr names>)] means you must specify all attribute values in the exact order defined in relation
- Partial tuple: specify a subset of the attribute values in the same order as the list of attributes [(<attr names>)]

SQL Modification: INSERT (3)

Inserting a tuple using SELECT command

Syntax:

INSERT INTO <table name>[(<attr names>)] (<**SELECT** subquery>)

- Multiple tuples may be added dependent on the **SELECT** subquery relation

SQL Example: INSERT

- Complete tuple:

```
INSERT INTO employee VALUES ('Joyce', 'C', 'Ho',  
'111223333', '1985-02-05', '400 Dowman Drive,  
Atlanta, GA', 'F', '150000', '987654321', 5);
```

- Partial tuple:

```
INSERT INTO employee(fname, lname, ssn) VALUES  
( 'Joyce', 'Ho', '111223333');
```

SQL Example: INSERT w/ SELECT

Suppose we want a new table that has the name, number of employees, and total salaries for each department. We first create the table then load it with the information from the database.

```
CREATE TABLE dept_info  
( dept_name  VARCHAR(10),  
  no_of_emps INT,  
  tot_salary  INT);
```

```
INSERT INTO dept_info  
  (SELECT      dname, count(*), sum(salary)  
   FROM        department, employee  
   WHERE       dnumber = dno  
   GROUP BY dname);
```

MySQL: Bulk Import

- All respectable RDBMS provide utilities to import data from text files
 - Syntax for uploading data will vary based on vendor
- MySQL allows the LOAD DATA INFILE (<http://dev.mysql.com/doc/refman/5.7/en/load-data.html>)
 - For a pipe-delimited file (| separates each column):
LOAD DATA LOCAL INFILE <filename>
{REPLACE | IGNORE} INTO TABLE <table name>
FIELDS TERMINATED BY '|';

SQL Modification: DELETE

- Remove tuples from a relation
- Syntax:
DELETE FROM <relation>
WHERE <condition>;
- Be careful! All tuples that satisfy the condition clause are deleted
- Tuples are deleted from only one table at a time unless CASCADE is specified on a referential integrity constraint
- What happens if we don't specify a WHERE clause?

SQL Example: DELETE

Delete all employees with the last name Brown

```
DELETE FROM employee  
WHERE lname = 'Brown';
```

SQL Example: DELETE (2)

Delete all employees from the 'Research' department who have more than 2 dependents

```
DELETE FROM employee
WHERE dno IN (SELECT dnumber
               FROM department
               WHERE dname = 'Research')
AND ssn IN (SELECT essn
            FROM dependent
            GROUP BY essn
            HAVING COUNT(name) > 2);
```


SQL Modification: UPDATE

- Modify/change certain attributes in certain tuples of a relation
- Syntax:
UPDATE <relation>
SET <list of attribute assignments>
WHERE <condition>;
- **UPDATE** command modifies tuples in the same relation

SQL Example: UPDATE

Change the location and controlling department number of project 10 to 'Bellaire' and 5, respectively.

```
UPDATE project  
SET      plocation = 'Bellaire', dnum = 5  
WHERE pnumber = 10;
```

SQL Example: UPDATE (2)

Give all employees in the 'Research' department a 10% raise

```
UPDATE employee
SET      salary = salary * 1.1
WHERE   dno IN (SELECT dnumber
                FROM   department
                WHERE  dname = 'Research');
```


- Reference to **salary** attribute on the right of = refers to the salary value before modification
- Reference to **salary** attribute on the left of = refers to salary value after modification

SQL: VIEW

- A view is a virtual table, a relation that is defined in terms of the contents of other tables and views
- A view does *not* exist in the physical form
- In contrast, a relation whose value is really in the database is called a base table
- Syntax:
CREATE VIEW <name> AS <query>;

SQL: View & Logical Data Independence

- Recall Logical Data Independence (class on Database Concepts)
 - Ability to present the stored information in a different way to different users
- View can be adapted to the need of the user
- If conceptual schema changes, only the SELECT query needed to construct view needs to change



SQL Example: VIEW

- Suppose an administrator maintains a list of activities of all employees which contains the following information:
fname, lname, project_name, hours_worked
- Regular SELECT query:
SELECT fname, lname, pname, hours
FROM employee, works-on, project
WHERE ssn = essn AND pno = pnumber;
- Create VIEW for the admin:
CREATE VIEW emp_activity
AS (SELECT fname, lname, pname, hours
FROM employee, works-on, project
WHERE ssn = essn AND pno = pnumber);

SQL: VIEW Advantages

- View can be used in queries like an ordinary relation
 - When a view is used in a SELECT query, the virtual relation is computed first
- Simplify complex queries by hiding them from the end-user and applications
- Limit data access to specific users (expose only non-sensitive data) and provides extra security for read/write access
- Enables backward compatibility - changes to database won't affect changes to other applications

SQL: VIEW Disadvantages

- Querying data from database view can be slow (since view is computed each time)
- Tables dependency - updates to the underlying tables will force changes to the view itself to make it work properly
- Most data manipulation statements (INSERT, DELETE, UPDATE) are not possible on the view

SQL Data Update & View: Recap

- Query
 - Temporal Relation / WITH
 - JOIN
- SQL Modification
 - INSERT
 - DELETE
 - UPDATE
- SQL Views



More SQL Practice

Find the fname, lname of employees with more than 2 dependents and work on all projects controlled by department #1

More SQL Practice

Find the fname, lname of employees with more than 2 dependents and work on all projects controlled by department #1

SELECT fname, lname
FROM employee

First formulate in words

WHERE <employee has more than 2 dependents>

AND <works on all projects controlled by dept #1>;

conquer each subquery separately

More SQL Practice: Subquery #1

Find employee that has more than 2 dependents

```
SELECT essn  
FROM dependent  
GROUP BY essn  
HAVING COUNT(name) > 2
```

More SQL Practice: Subquery #2

Find employees that works on all projects controlled by department #1 - set difference technique

```
SELECT ssn
FROM   employee e
WHERE  NOT EXISTS
    < set of projects controlled by department #1 >
    - <set of projects worked on by e.ssn>;
```

More SQL Practice: Subquery #2

Find employees that works on all projects controlled by department #1 - set difference technique

```
SELECT ssn
FROM   employee e
WHERE NOT EXISTS (SELECT pnumber
                  FROM   project
                  WHERE pnumber IN
                        ( SELECT pnumber
                          FROM   project
                          WHERE dnum = 1)
                  AND pnumber NOT IN
                        ( SELECT pno
                          FROM   works_on
                          WHERE essn = e.ssn));
```

More SQL Practice: Putting it Together

Find the fname, lname of employees with more than 2 dependents and work on all projects controlled by department #1

```
SELECT fname, lname
FROM   employee
WHERE  ssn IN ( SELECT essn
                FROM   dependent
                GROUP BY essn
                HAVING COUNT(name) > 2)
AND    ssn IN ( SELECT ssn
                FROM   employee e
                WHERE NOT EXISTS (SELECT pnumber
                                FROM   project
                                WHERE pnumber IN
                                      (SELECT pnumber
                                       FROM   project
                                       WHERE dnum = 1)
                                AND pnumber NOT IN
                                      ( SELECT pno
                                       FROM   works_on
                                       WHERE essn = e.ssn))));
```

More SQL Practice (2)

Find the department name, and the number of employees in that department that earns more than 40K for departments with at least 2 employees

More SQL Practice (2)

Find the department name, and the number of employees in that department that earns more than 40K for departments with at least 2 employees

```
SELECT    dname, COUNT(ssn)
FROM      department, employee
WHERE     dnumber = dno
        AND salary > 40000
GROUP BY  dname
HAVING    COUNT(ssn) >= 2;
```

What is wrong with this solution?

More SQL Practice (2): Solution

Find the department name, and the number of employees in that department that earns more than 40K for departments with at least 2 employees

```
SELECT    dname, COUNT(ssn)
FROM      department, employee
WHERE     dnumber = dno
        AND salary > 40000
        AND dno IN ( SELECT    dno
                        FROM      employee
                        GROUP BY dno
                        HAVING     COUNT(ssn) >= 2 )
GROUP BY dname;
```

More SQL Practice (3)

Find fname, name of employees who work on 2 or more projects together with John Smith

More SQL Practice (4)

Find departments who have 2 or more employees working on all projects controlled by 'Research' department